

IE411: Operating Systems

Deadlock bugs

Deadlocks

- Deadlock: No progress can be made because two or more threads are waiting for the other to take some action and thus neither ever does
- Could arise when we need to coordinate access to more than one shared resource
 - e.g., acquiring multiple locks can result in deadlock

Example

Thread 1:

lock(L1);

lock(L2);

Thread 2:

lock(L2);

lock(L1);

- If this code runs, deadlock does not necessarily occur
- It may occur, if, for example
 - Thread 1 grabs lock L1 and then a context switch occurs to Thread 2
 - At that point, Thread 2 grabs L2, and tries to acquire L1

Conditions for deadlock

- Four conditions need to hold for a deadlock to occur
 - ① Mutual exclusion: threads claim exclusive control of resources that they require
 - ② Hold-and-wait: threads hold resources allocated to them (e.g., locks that they have already acquired) while waiting for additional resources (e.g., locks)

Conditions for deadlock

- Four conditions need to hold for a deadlock to occur
 - ① Mutual exclusion: threads claim exclusive control of resources that they require
 - ② Hold-and-wait: threads hold resources allocated to them (e.g., locks that they have already acquired) while waiting for additional resources (e.g., locks)
 - ③ No preemption: resources (e.g., locks) cannot be forcibly removed from threads that are holding them

Conditions for deadlock

- Four conditions need to hold for a deadlock to occur
 - ① Mutual exclusion: threads claim exclusive control of resources that they require
 - ② Hold-and-wait: threads hold resources allocated to them (e.g., locks that they have already acquired) while waiting for additional resources (e.g., locks)
 - ③ No preemption: resources (e.g., locks) cannot be forcibly removed from threads that are holding them
 - ④ Circular wait: there exists a circular chain of threads such that each thread holds one or more resources (e.g., locks) that are being requested by the next thread in the chain

Deadlock prevention

- (mutual exclusion) and (hold-and-wait) and (no-preemption) and (circular wait)
- If any of these four conditions are not met deadlock cannot occur

Deadlock prevention

- (mutual exclusion) and (hold-and-wait) and (no-preemption) and (circular wait)
- If any of these four conditions are not met deadlock cannot occur
- Deadlock prevention: prevent one of the above conditions from arising
 - \sim (mutual exclusion) or \sim (hold-and-wait) or \sim (no preemption) or \sim (circular wait) \rightarrow no deadlock

Preventing circular wait

- Avoid circular wait by acquiring locks in the same order
 - called a locking hierarchy
- Example
 - Assume two threads need locks L1 and L2
 - Use a locking hierarchy, e.g., $L1 \rightarrow L2$

Thread 1:

`lock(L1);`

`lock(L2);`

Thread 2:

`lock(L1);`

`lock(L2);`

Preventing circular wait

Thread 1:

`lock(L1);`

`lock(L2);`

Thread 2:

`lock(L1);`

`lock(L2);`

- Only one thread will be able to acquire L1, the other thread has to block
- No deadlock possible, because the blocking thread did not acquire L2 before if it obeys the locking hierarchy
- The blocking thread can continue only after the locks have been released

Enforce lock ordering

- Example: `set_t *intersect(set_t *s1, set_t *s2)`
- Locks `s1->lock` and `s2->lock` for the sets `s1` and `s2` need to be acquired
- Assume `intersect()` acquires the said locks in some fixed order
 - `s1->lock` before `s2->lock` (or `s2->lock` before `s1->lock`)
- Deadlock possible?

Enforce lock ordering

- Example: `set_t *intersect(set_t *s1, set_t *s2)`
- Locks `s1->lock` and `s2->lock` for the sets `s1` and `s2` need to be acquired
- Assume `intersect()` acquires the said locks in some fixed order
 - `s1->lock` before `s2->lock` (or `s2->lock` before `s1->lock`)
- Deadlock possible?
 - one thread calls `intersect(s1, s2)`
 - another thread calls `intersect(s2, s1)`

Enforce lock ordering

- One possible solution is to acquire the locks in the order of their virtual addresses

```
if (&s1->lock > &s2->lock) {  
    // grab locks in high-to-low address order  
    lock(s1->lock);  
    lock(s2->lock);  
} else {  
    lock(s2->lock);  
    lock(s1->lock);  
}
```

Preventing hold-and-wait

- Avoid hold-and-wait by acquiring all locks at once atomically, say by acquiring a master lock first

```
1: lock ( prevention );  
2: lock ( L1 );  
3: lock ( L2 );  
4: ...  
5: unlock ( prevention );
```

Preventing hold-and-wait

- Avoid hold-and-wait by acquiring all locks at once atomically, say by acquiring a master lock first

```
1: lock (prevention);  
2: lock (L1);  
3: lock (L2);  
4: ...  
5: unlock (prevention);
```

- Guarantees that no untimely thread switch can occur in the midst of lock acquisition
- But this method may reduce concurrent execution and performance gains

Preventing no preemption

- Strategy: If a thread can't get what it wants, release what it holds
- trylock()
 - grabs the lock if it is available, otherwise returns -1

```
1: top:
2:     lock(L1);
3:     if (trylock(L2) == -1) {
4:         unlock(L1);
5:         goto top;
6:     }
```


Preventing no preemption

<pre>1: top: 2: lock(L1); 3: if (trylock(L2) == -1) { 4: unlock(L1); 5: goto top; 6: }</pre>	<pre>top: lock(L2); if (trylock(L1) == -1) { unlock(L2); goto top; }</pre>
--	--

- Another thread could follow the same protocol but grab locks in the other order (L2 then L1)
 - still be deadlock free

New issue: Livelock

- It can happen that both threads repeatedly attempt this sequence and repeatedly fail to acquire both locks
- Livelock: the state of the threads is constantly changing (so not a deadlock), but there is no progress

New issue: Livelock

- It can happen that both threads repeatedly attempt this sequence and repeatedly fail to acquire both locks
- Livelock: the state of the threads is constantly changing (so not a deadlock), but there is no progress
- solution: add a random delay looping back and trying the entire thing over again

Preventing mutual exclusion

- Avoid the need for mutual exclusion at all
- But how, the code we wish to run does have critical sections
- Use lock-free approaches to code data structures without explicit locking

Compare and swap

Yet another atomic instruction (pseudo C code):

```
int CompareAndSwap(int *address, int expected, int new) {  
    if (*address == expected) {  
        *address = new;  
        return 1;  
    }  
    return 0;  
}
```

Atomic increase

- Goal: atomically increment a value by a certain amount

```
void AtomicIncrement(int *value, int amount) {  
    do {  
        int old = *value;  
    } while (CompareAndSwap(value, old, old+amount) == 0);  
}
```

- no lock is required and no deadlock can arise (livelock is possible)

Other solutions to deadlock

- Deadlock avoidance
 - requires global knowledge of which locks various threads might grab during their execution, and subsequently schedule threads in a way as to guarantee no deadlock can occur
 - Banker's algorithm is a well-known deadlock avoidance algorithm
- Detect and recover (reboot or kill deadlock threads)

Next time

- Lock-free linked lists!

Exercie: Spinlock using compare-and-swap

```
typedef struct {  
    int flag;  
} lock_t;  
  
void init(lock_t *lock) {  
    // 0 indicates that lock is available ,  
    // 1 that it is held  
    lock->flag = 0;  
}  
  
void lock(lock_t *lock) {  
    while (CAS(?, ?, ?) == ?)  
        ; // spin-wait (do nothing)  
}  
  
void unlock(lock_t *lock) {  
    lock->flag = 0;  
}
```