

# IE 411: Operating Systems

## Non-blocking linked lists

# A sorted linked list

```
struct Node {  
    int value;  
    Node* next;  
};
```

```
struct List {  
    Node* head;  
};
```

```
void insert(List* list, int value) {  
  
    Node* n = new Node;  
    n->value = value;  
  
    // assume case of inserting before head of  
    // of list is handled here (to keep slide simple)  
  
    Node* prev = list->head;  
    Node* cur = list->head->next;  
  
    while (cur) {  
        if (cur->value > value)  
            break;  
  
        prev = cur;  
        cur = cur->next;  
    }  
  
    n->next = cur;  
    prev->next = n;  
}
```

## What can go wrong if multiple threads operate on the linked list simultaneously?


```
void delete(List* list, int value) {  
  
    // assume case of deleting first element is  
    // handled here (to keep slide simple)  
  
    Node* prev = list->head;  
    Node* cur = list->head->next;  
  
    while (cur) {  
        if (cur->value == value) {  
            prev->next = cur->next;  
            delete cur;  
            return;  
        }  
  
        prev = cur;  
        cur = cur->next;  
    }  
}
```

## Example: concurrent insertion



- Thread 1 attempts to insert 6
- Thread 2 attempts to insert 7



**Thread 1:**




# Example: concurrent insertion



Thread 1 and thread 2 both compute same prev and cur.  
Result: one of the insertions gets lost!

**Result: (assuming thread 1 updates prev->next before thread 2)**



# Solution 1: global locking

```
struct Node {  
    int value;  
    Node* next;  
};
```

```
struct List {  
    Node* head;  
    Lock lock;
```

← Per-list lock

```
void insert(List* list, int value) {  
  
    Node* n = new Node;  
    n->value = value;  
  
    lock(list->lock);  
  
    // assume case of inserting before head of  
    // of list is handled here (to keep slide simple)  
  
    Node* prev = list->head;  
    Node* cur = list->head->next;  
  
    while (cur) {  
        if (cur->value > value)  
            break;  
  
        prev = cur;  
        cur = cur->next;  
    }  
    n->next = cur;  
    prev->next = n;  
    unlock(list->lock);  
}
```

```
void delete(List* list, int value) {  
  
    lock(list->lock);  
  
    // assume case of deleting first element is  
    // handled here (to keep slide simple)  
  
    Node* prev = list->head;  
    Node* cur = list->head->next;  
  
    while (cur) {  
        if (cur->value == value) {  
            prev->next = cur->next;  
            delete cur;  
            unlock(list->lock);  
            return;  
        }  
  
        prev = cur;  
        cur = cur->next;  
    }  
    unlock(list->lock);  
}
```

# Single global lock

- Advantages
  - simple to implement

# Single global lock

- Advantages
  - simple to implement
- Disadvantages?

# Single global lock

- Advantages
  - simple to implement
- Disadvantages?
  - Operations on the data structure are serialized



# Single global lock

- Advantages
  - simple to implement
- Disadvantages?
  - Operations on the data structure are serialized
  - May limit application performance

# Lock-free algorithms

- protecting DS (e.g. BST, linked list) with a single lock is pessimistic as it assumes conflicts will occur
- a lockless algorithm is optimistic as it assumes conflicts unlikely to occur and, when they are detected, they are resolved

# Lock-free algorithms

- protecting DS (e.g. BST, linked list) with a single lock is pessimistic as it assumes conflicts will occur
- a lockless algorithm is optimistic as it assumes conflicts unlikely to occur and, when they are detected, they are resolved
- Advantages compared to locking?

# Lock-free algorithms

- protecting DS (e.g. BST, linked list) with a single lock is pessimistic as it assumes conflicts will occur
- a lockless algorithm is optimistic as it assumes conflicts unlikely to occur and, when they are detected, they are resolved
- Advantages compared to locking?
  - allows concurrency while there are no conflicts **which hopefully is so most of the time**


# Atomic Compare-and-Swap (CAS)

```
bool CAS(  
    memory location L,  
    expected value V at L,  
    desired new value V1 at L  
);
```

If (the expected value V at memory location L == the current value at L), CAS succeeds by storing the the desired value V1 at L and returns TRUE.


# Adding nodes

- Ex: use CAS to add nodes 15 and 35



# Adding nodes

- Ex: use CAS to add nodes 15 and 35




- search for insertion point, initialise next pointer and then execute with correct parameters to insert node into list

```
CAS(&a->next, b, c);    // add node c between a and b  
CAS(&d->next, e, f);    // add node f between d and e
```

- disjoint-access parallelism

# Adding nodes

- if 2 threads try to add nodes at the same position



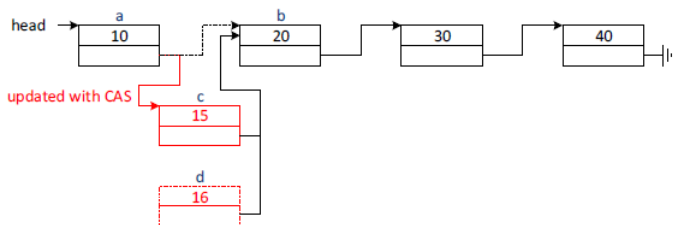
```
CAS(&a->next, b, c); // first CAS executed will succeed..  
CAS(&a->next, b, d); // and thus second CAS executed will FAIL
```

- first CAS executed succeeds, second will fail as `a->next != b`



# Adding nodes

- if 2 threads try to add nodes at the same position




```
CAS(&a->next, b, c); // first CAS executed will succeed..  
CAS(&a->next, b, d); // and thus second CAS executed will FAIL
```

- first CAS executed succeeds, second will fail as  $a \rightarrow next \neq b$
- RETRY on failure, which means searching for insertion point AGAIN and, if key not found, set up and re-execute CAS

# Removing nodes

- search for node and then execute CAS with correct parameters to remove node from list
- consider 2 threads removing non-adjacent nodes



```
CAS(&a->next, b, c); // remove node b (20)
```

```
CAS(&c->next, d, 0); // remove node d (40)
```

- disjoint access parallelism

# Removing nodes

- if two threads try to remove the same node




```
CAS(&a->next, b, c);
```

```
CAS(&a->next, b, c);
```

# Removing nodes


- if two threads try to remove the same node



```
CAS(&a->next, b, c);
```

```
CAS(&a->next, b, c);
```

- assume first CAS executed succeeds



- then second CAS executed fails as  $a \rightarrow next \neq b$

# Removing nodes


- if two threads try to remove the same node



```
CAS(&a->next, b, c);
```

```
CAS(&a->next, b, c);
```


- assume first CAS executed succeeds




- then second CAS executed fails as  $a \rightarrow \text{next} \neq b$
- RETRY on failure, which means searching AGAIN for node (which may not be found)

# What doesn't work...

- consider removing node 20 and adding node 25 concurrently




```
CAS(&a->next, b, c); // remove 20  
CAS(&b->next, c, d); // add 25
```




NB: same result if CAS instructions  
executed in reverse order

# What doesn't work...

- consider removing node 20 and adding node 25 concurrently



```
CAS(&a->next, b, c); // remove 20  
CAS(&b->next, c, d); // add 25
```




NB: same result if CAS instructions  
executed in reverse order

- NOT what was intended!


# What doesn't work...

- imagine deleting adjacent nodes



```
CAS(&a->next, b, c); // remove 20
```

```
CAS(&b->next, c, d); // remove 30
```




- AGAIN NOT what was intended!




# ABA Problem

- imagine insertion point found, BUT before `CAS(&a->next, b, c)` is executed, thread is pre-empted




- another thread then removes b from list




# ABA Problem

- imagine insertion point found, BUT before `CAS(&a->next, b, c)` is executed, thread is pre-empted




- another thread then removes b from list




- if thread adding 15 resumes execution, the CAS fails which is OK in this case

# ABA Problem

- imagine insertion point found, BUT before `CAS(&a->next, b, c)` is executed, thread is pre-empted




- another thread then removes b from list



- if thread adding 15 resumes execution, the CAS fails which is OK in this case
- BUT what bad thing can happen?

# ABA Problem

- if the memory used by b is reused, for example by a thread adding key 12 to the list before thread adding 15 resumes ...
- when the thread adding 15 to list resumes, its CAS will succeed and 15 will be added into the list at the wrong position



# ABA Problem

- avoid the ABA problem by not reusing nodes:
  - nodes cannot be reused if any thread has or can get a pointer to the node

# ABA Problem

- avoid the ABA problem by not reusing nodes:
  - nodes cannot be reused if any thread has or can get a pointer to the node
- Disadvantages?

# ABA Problem

- avoid the ABA problem by not reusing nodes:
  - nodes cannot be reused if any thread has or can get a pointer to the node
- Disadvantages?
  - will quickly run out of memory