



IE411: Operating Systems

Semaphores

Synchronization constructs

- Previously we looked at locks and condition variables
- We can combine these two concepts to create a new type of synchronization primitive called semaphore

Equivalence



Semaphore data structure

```
typedef struct {  
    int count;  
    queue_t waiting;  
} semaphore_t;
```

- an integer value
- a queue of threads waiting on the semaphore

Semaphore operations - sem_wait()

- `sem_wait()` decrements the semaphore's value
 - also called P() (Dijkstra - Dutch for prolaag)
- if the resulting value is negative, the calling thread is blocked
 - a blocked thread can be woken up when the semaphore value is incremented

```
void sem_wait(semaphore_t *s)
{
    s->count--;
    if (s->count < 0) {
        add calling thread to s->queue;
        put calling thread to sleep;
    }
}
```

Semaphore operations - `sem_post()`

- `sem_post()` increments the semaphore's value
 - also called V() (Dijkstra - Dutch for verhoog)
- incrementing the value when it is negative causes one of the waiter threads to become runnable

```
void sem_post(semaphore_t *s)
{
    s->count++;
    if (s->count <= 0) {
        remove one thread from s->queue;
        wake up this thread;
    }
}
```

Semaphore nuance

- `sem_wait()` and `sem_post()` are made atomic
- how negative the semaphore's value is represents the number of threads waiting

Application - Mutual Exclusion

- like a lock
- integer value represents a 1/0 value to lock and unlock a critical section

Val	Thread 0	State	Thread 1	State
1		Run		Ready
1	call sem_wait()	Run		Ready
0	sem_wait() returns	Run		Ready
0	(crit sect begin)	Run		Ready
0	<i>Interrupt; Switch→T1</i>	Ready		Run
0		Ready	call sem_wait()	Run
-1		Ready	decr sem	Run
-1		Ready	(sem<0)→sleep	Sleep
-1		Run	<i>Switch→T0</i>	Sleep
-1	(crit sect end)	Run		Sleep
-1	call sem_post()	Run		Sleep
0	incr sem	Run		Sleep
0	wake(T1)	Run		Ready
0	sem_post() returns	Run		Ready
0	<i>Interrupt; Switch→T1</i>	Ready		Run
0		Ready	sem_wait() returns	Run
0		Ready	(crit sect)	Run
0		Ready	call sem_post()	Run
1		Ready	sem_post() returns	Run

Figure 31.5: Thread Trace: Two Threads Using A Semaphore

Application - Ordering

- Like a condition variable
- If we choose the correct starting value, we can ensure some simple ordering scenarios
- $X = ?$

```
1  sem_t s;  
2  
3  void *child(void *arg) {  
4      printf("child\n");  
5      sem_post(&s); // signal here: child is done  
6      return NULL;  
7  }  
8  
9  int main(int argc, char *argv[]) {  
10     sem_init(&s, 0, X); // what should X be?  
11     printf("parent: begin\n");  
12     pthread_t c;  
13     Pthread_create(&c, NULL, child, NULL);  
14     sem_wait(&s); // wait here for child  
15     printf("parent: end\n");  
16     return 0;  
17 }
```

Figure 31.6: A Parent Waiting For Its Child

Application - Ordering

```

1 sem_t s;
2
3 void *child(void *arg) {
4     printf("child\n");
5     sem_post(&s); // signal here: child is done
6     return NULL;
7 }
8
9 int main(int argc, char *argv[]) {
10     sem_init(&s, 0, X); // what should X be?
11     printf("parent: begin\n");
12     pthread_t c;
13     Pthread_create(&c, NULL, child, NULL);
14     sem_wait(&s); // wait here for child
15     printf("parent: end\n");
16     return 0;
17 }

```

Figure 31.6: A Parent Waiting For Its Child

Val	Parent	State	Child	State
0	create (Child)	Run	(Child exists, can run)	Ready
0	call sem_wait ()	Run		Ready
-1	decr sem	Run		Ready
-1	(sem<0)→sleep	Sleep		Ready
-1	Switch→Child	Sleep	child runs	Run
-1		Sleep	call sem_post ()	Run
0		Sleep	inc sem	Run
0		Ready	wake (Parent)	Run
0		Ready	sem_post () returns	Run
0		Ready	Interrupt→Parent	Ready
0	sem.wait () returns	Run		Ready

Figure 31.7: Thread Trace: Parent Waiting For Child (Case 1)

Application - Ordering

```

1 sem_t s;
2
3 void *child(void *arg) {
4     printf("child\n");
5     sem_post(&s); // signal here: child is done
6     return NULL;
7 }
8
9 int main(int argc, char *argv[]) {
10     sem_init(&s, 0, X); // what should X be?
11     printf("parent: begin\n");
12     pthread_t c;
13     pthread_create(&c, NULL, child, NULL);
14     sem_wait(&s); // wait here for child
15     printf("parent: end\n");
16     return 0;
17 }

```

Figure 31.6: A Parent Waiting For Its Child

Val	Parent	State	Child	State
0	create (Child)	Run	(Child exists; can run)	Ready
0	<i>Interrupt</i> → Child	Ready	child runs	Run
0		Ready	call sem_post ()	Run
1		Ready	inc sem	Run
1		Ready	wake (nobody)	Run
1		Ready	sem_post () returns	Run
1	parent runs	Run	<i>Interrupt</i> → Parent	Ready
1	call sem_wait ()	Run		Ready
0	decrement sem	Run		Ready
0	(sem ≥ 0) → awake	Run		Ready
0	sem_wait () returns	Run		Ready

Figure 31.8: Thread Trace: Parent Waiting For Child (Case 2)

Producer-Consumer Problem

- The bounded-buffer producer-consumer problem assumes there is a buffer of size N
- The producer puts items to the buffer area
- The consumer consumes items from the buffer
- The producer and consumer execute concurrently

Producer-Consumer Model

- Shared data

```
sem_t full , empty;
```

- Initially

```
full = 0;           /* The number of full buffers */  
empty = MAX;        /* The number of empty buffers */
```

First Attempt: $MAX = 1$

```
1  sem_t empty;
2  sem_t full;
3
4  void *producer(void *arg) {
5      int i;
6      for (i = 0; i < loops; i++) {
7          sem_wait(&empty);          // line P1
8          put(i);                    // line P2
9          sem_post(&full);           // line P3
10     }
11 }
12
13 void *consumer(void *arg) {
14     int i, tmp = 0;
15     while (tmp != -1) {
16         sem_wait(&full);            // line C1
17         tmp = get();                // line C2
18         sem_post(&empty);           // line C3
19         printf("%d\n", tmp);
20     }
21 }
22
23 int main(int argc, char *argv[]) {
24     // ...
25     sem_init(&empty, 0, MAX); // MAX buffers are empty to begin with...
26     sem_init(&full, 0, 0);    // ... and 0 are full
27     // ...
28 }
```

```
1  int buffer[MAX];
2  int fill = 0;
3  int use = 0;
4
5  void put(int value) {
6      buffer[fill] = value;
7      fill = (fill + 1) % MAX;
8  }
9
10 int get() {
11     int tmp = buffer[use];
12     use = (use + 1) % MAX;
13     return tmp;
14 }
```

Put and Get routines

First Attempt: MAX = 10?

```
1  sem_t empty;
2  sem_t full;
3
4  void *producer(void *arg) {
5      int i;
6      for (i = 0; i < loops; i++) {
7          sem_wait(&empty);          // line P1
8          put(i);                    // line P2
9          sem_post(&full);           // line P3
10     }
11 }
12
13 void *consumer(void *arg) {
14     int i, tmp = 0;
15     while (tmp != -1) {
16         sem_wait(&full);             // line C1
17         tmp = get();                 // line C2
18         sem_post(&empty);           // line C3
19         printf("%d\n", tmp);
20     }
21 }
22
23 int main(int argc, char *argv[]) {
24     // ...
25     sem_init(&empty, 0, MAX); // MAX buffers are empty to begin with...
26     sem_init(&full, 0, 0);    // ... and 0 are full
27     // ...
28 }
```

```
1  int buffer[MAX];
2  int fill = 0;
3  int use = 0;
4
5  void put(int value) {
6      buffer[fill] = value;
7      fill = (fill + 1) % MAX;
8  }
9
10 int get() {
11     int tmp = buffer[use];
12     use = (use + 1) % MAX;
13     return tmp;
14 }
```

Put and Get routines


First Attempt: $MAX = 10$?

fill = 0

empty = 10


Producer 0: **Running**

```
void *producer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        sem_wait(&empty);  
        put(i);  
        sem_post(&full);  
    }  
}
```



Producer 1: **Runnable**

```
void *producer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        sem_wait(&empty);  
        put(i);  
        sem_post(&full);  
    }  
}
```




First Attempt: $MAX = 10$?

fill = 0


empty = 9

Producer 0: Running




```
void *producer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        sem_wait(&empty);  
        put(i);  
        sem_post(&full);  
    }  
}
```

Producer 1: Runnable



```
void *producer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        sem_wait(&empty);  
        put(i);  
        sem_post(&full);  
    }  
}
```



```
void put(int value) {  
    buffer[fill] = value;  
    fill = (fill + 1) % MAX;  
}
```


First Attempt: MAX = 10?

fill = 0

empty = 9


Producer 0: **Running**

```
void *producer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        sem_wait(&empty);  
        put(i);  
        sem_post(&full);  
    }  
}
```




Producer 1: **Runnable**

```
void *producer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        sem_wait(&empty);  
        put(i);  
        sem_post(&full);  
    }  
}
```



```
void put(int value) {  
    buffer[fill] = value;  
    Interrupted ...  
    fill = (fill + 1) % MAX;  
}
```




First Attempt: $MAX = 10$?

fill = 0

empty = 9


Producer 0: **Sleeping**

```
void *producer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        sem_wait(&empty);  
        put(i);  
        sem_post(&full);  
    }  
}
```




Producer 1: Runnable

```
void *producer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        sem_wait(&empty);  
        put(i);  
        sem_post(&full);  
    }  
}
```



```
void put(int value) {  
    buffer[fill] = value;  
    Interrupted ...  
    fill = (fill + 1) % MAX;  
}
```




First Attempt: $MAX = 10$?

fill = 0

empty = 9


Producer 0: Runnable



```
void *producer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        sem_wait(&empty);  
        put(i);  
        sem_post(&full);  
    }  
}
```

```
void put(int value) {  
    buffer[fill] = value;  
    Interrupted ...  
    fill = (fill + 1) % MAX;  
}
```

Producer 1: Running



```
void *producer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        sem_wait(&empty);  
        put(i);  
        sem_post(&full);  
    }  
}
```

First Attempt: MAX = 10?


fill = 0

Overwrite!

empty = 8

Producer 0: Runnable


```
void *producer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        sem_wait(&empty);  
        put(i);  
        sem_post(&full);  
    }  
}
```




```
void put(int value) {  
    buffer[fill] = value;  
    Interrupted ...  
    fill = (fill + 1) % MAX;  
}
```

Producer 1: Running

```
void *producer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        sem_wait(&empty);  
        put(i);  
        sem_post(&full);  
    }  
}
```



```
void put(int value) {  
     buffer[fill] = value;  
    fill = (fill + 1) % MAX;  
}
```

Producer-Consumer Model

- Shared data

```
sem_t full , empty , mutex;
```

- Initially

```
full = 0;           /* The number of full buffers */  
empty = MAX;        /* The number of empty buffers */  
mutex = 1;          /* Semaphore controlling the access  
                     to the buffer pool */
```

Add Mutual Exclusion

```
1  sem_t empty;
2  sem_t full;
3  sem_t mutex;
4
5  void *producer(void *arg) {
6      int i;
7      for (i = 0; i < loops; i++) {
8          sem_wait(&mutex);          // line p0 (NEW LINE)
9          sem_wait(&empty);          // line p1
10         put(i);                    // line p2
11         sem_post(&full);            // line p3
12         sem_post(&mutex);          // line p4 (NEW LINE)
13     }
14 }
15
16 void *consumer(void *arg) {
17     int i;
18     for (i = 0; i < loops; i++) {
19         sem_wait(&mutex);          // line c0 (NEW LINE)
20         sem_wait(&full);            // line c1
21         int tmp = get();           // line c2
22         sem_post(&empty);          // line c3
23         sem_post(&mutex);          // line c4 (NEW LINE)
24         printf("%d\n", tmp);
25     }
26 }
27
28 int main(int argc, char *argv[]) {
29     // ...
30     sem_init(&empty, 0, MAX); // MAX buffers are empty to begin with...
31     sem_init(&full, 0, 0);    // ... and 0 are full
32     sem_init(&mutex, 0, 1);   // mutex=1 because it is a lock (NEW LINE)
33     // ...
34 }
```

Add Mutual Exclusion


```
1  sem_t empty;
2  sem_t full;
3  sem_t mutex;
4
5  void *producer(void *arg) {
6      int i;
7      for (i = 0; i < loops; i++) {
8          sem_wait(&mutex);          // line p0 (NEW LINE)
9          sem_wait(&empty);          // line p1
10         put(i);                     // line p2
11         sem_post(&full);            // line p3
12         sem_post(&mutex);           // line p4 (NEW LINE)
13     }
14 }
15
16 void *consumer(void *arg) {
17     int i;
18     for (i = 0; i < loops; i++) {
19         sem_wait(&mutex);          // line c0 (NEW LINE)
20         sem_wait(&full);           // line c1
21         int tmp = get();           // line c2
22         sem_post(&empty);          // line c3
23         sem_post(&mutex);           // line c4 (NEW LINE)
24         printf("%d\n", tmp);
25     }
26 }
27
28 int main(int argc, char *argv[]) {
29     // ...
30     sem_init(&empty, 0, MAX); // MAX buffers are empty to begin with...
31     sem_init(&full, 0, 0);    // ... and 0 are full
32     sem_init(&mutex, 0, 1);   // mutex=1 because it is a lock (NEW LINE)
33     // ...
34 }
```

**What if consumer
gets to run first??**

Add Mutual Exclusion


mutex = 1
full = 0
empty = 10

Producer 0: Runnable



```
void *producer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        sem_wait(&mutex);  
        sem_wait(&empty);  
        put(i);  
        sem_post(&full);  
        sem_post(&mutex);  
    }  
}
```

Consumer 0: **Running**



```
void *consumer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        sem_wait(&mutex);  
        sem_wait(&full);  
        int tmp = get();  
        sem_post(&empty);  
        sem_post(&mutex);  
        printf("%d\n", tmp);  
    }  
}
```


Add Mutual Exclusion

mutex = 0

full = 0


empty = 10

Producer 0: Runnable



```
void *producer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        sem_wait(&mutex);  
        sem_wait(&empty);  
        put(i);  
        sem_post(&full);  
        sem_post(&mutex);  
    }  
}
```

Consumer 0: **Running**



```
void *consumer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        sem_wait(&mutex);  
        sem_wait(&full);  
        int tmp = get();  
        sem_post(&empty);  
        sem_post(&mutex);  
        printf("%d\n", tmp);  
    }  
}
```

Consumer 0 is waiting for
full to be greater than or
equal to 0


Add Mutual Exclusion

mutex = -1

full = -1


empty = 10

Producer 0: Running



```
void *producer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        sem_wait(&mutex);  
        sem_wait(&empty);  
        put(i);  
        sem_post(&full);  
        sem_post(&mutex);  
    }  
}
```

Consumer 0: Runnable



```
void *consumer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        sem_wait(&mutex);  
        sem_wait(&full);  
        int tmp = get();  
        sem_post(&empty);  
        sem_post(&mutex);  
        printf("%d\n", tmp);  
    }  
}
```

Consumer 0 is **waiting** for
full to be greater than or
equal to 0

Add Mutual Exclusion


Deadlock!!

mutex = -1

full = -1


empty = 10

Producer 0: **Running**



```
void *producer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        sem_wait(&mutex);  
        sem_wait(&empty);  
        put(i);  
        sem_post(&full);  
        sem_post(&mutex);  
    }  
}
```

Consumer 0: **Runnable**




```
void *consumer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        sem_wait(&mutex);  
        sem_wait(&full);  
        int tmp = get();  
        sem_post(&empty);  
        sem_post(&mutex);  
        printf("%d\n", tmp);  
    }  
}
```

Producer 0 **gets stuck** at
acquiring **mutex** which has
been locked by Consumer 0!

Consumer 0 is **waiting** for
full to be greater than or
equal to 0

Deadlocks

When every thread in a set of threads is waiting for an event that can be caused only by another thread in the set



A typical deadlock
dependency graph

Correct Mutual Exclusion

```
1  sem_t empty;
2  sem_t full;
3  sem_t mutex;
4
5  void *producer(void *arg) {
6      int i;
7      for (i = 0; i < loops; i++) {
8          sem_wait(&empty);          // line p1
9          sem_wait(&mutex);          // line p1.5 (MOVED MUTEX HERE...)
10         put(i);                    // line p2
11         sem_post(&mutex);          // line p2.5 (... AND HERE)
12         sem_post(&full);           // line p3
13     }
14 }
15
16 void *consumer(void *arg) {
17     int i;
18     for (i = 0; i < loops; i++) {
19         sem_wait(&full);            // line c1
20         sem_wait(&mutex);          // line c1.5 (MOVED MUTEX HERE...)
21         int tmp = get();            // line c2
22         sem_post(&mutex);          // line c2.5 (... AND HERE)
23         sem_post(&empty);          // line c3
24         printf("%d\n", tmp);
25     }
26 }
27
28 int main(int argc, char *argv[]) {
29     // ...
30     sem_init(&empty, 0, MAX); // MAX buffers are empty to begin with...
31     sem_init(&full, 0, 0);    // ... and 0 are full
32     sem_init(&mutex, 0, 1);   // mutex=1 because it is a lock
33     // ...
34 }
```

**Mutex wraps
just around
critical section!**

**Mutex wraps
just around
critical section!**

- Single shared object
- Want to allow any number of threads to read simultaneously
- But, only one thread should be able to write to the object at a time
 - And, not interfere with any readers ...

Readers/Writers

- readers share:
 - semaphore wrt; // initialized to 1
 - int readcount; // initialized to 0
- writers also share semaphore wrt

Writer:

```
while (1) {  
    P(wrt);  
    // writing  
    V(wrt);  
}
```

Reader:

```
while (1) {  
    readcount++;  
    if (readcount == 1) {  
        P(wrt);  
    }  
    // reading  
    readcount--;  
    if (readcount == 0) {  
        V(wrt);  
    }  
}
```

- Seems simple, but this code is broken ...

Readers/Writers

Writer:

```
while (1) {  
    P(wrt);  
    // writing  
    V(wrt);  
}
```

Reader:

```
while (1) {  
    readcount++;  
    if (readcount == 1) {  
        P(wrt);  
    }  
    // reading  
    readcount--;  
    if (readcount == 0) {  
        V(wrt);  
    }  
}
```

*What can
happen if we
context switch
here?*

Writer:

```
while (1) {  
    P(wrt);  
    // writing  
    V(wrt);  
}
```

Reader:

```
while (1) {  
    readcount++;  
    if (readcount == 1) {  
        P(wrt);  
    }  
    // reading  
    readcount--;  
    if (readcount == 0) {  
        V(wrt);  
    }  
}
```

*Another Reader()
could start and
"readcount==1"
never happens!*

Writer:

```
while (1) {  
    P(wrt);  
    // writing  
    V(wrt);  
}
```

Reader:

```
while (1) {  
    readcount++;  
    if (readcount == 1) {  
        P(wrt);  
    }  
    // reading  
    readcount--;  
    if (readcount == 0) {  
        V(wrt);  
    }  
}
```

*What can
happen if we
context switch
here?*

Readers/Writers

A Writer() could start, P the semaphore first, then subsequent Reader() threads would be able to get past the semaphore (since "readcount != 1")

Writer.

```
while (1) {  
    P(wrt);  
    // writing  
    V(wrt);  
}
```

Reader.

```
while (1) {  
    readcount++;  
    if (readcount == 1) {  
        P(wrt);  
    }  
    // reading  
    readcount--;  
    if (readcount == 0) {  
        V(wrt);  
    }  
}
```

Readers/Writers fixed

- Problem: Multiple Readers are accessing readcount
- Solution: Make “increment, test, P” and “decrement, test, V” both atomic - using a mutex

Writer:

```
while (1) {  
    P(wrt);  
    // writing  
    V(wrt);  
}
```

Reader:

```
while (1) {  
    P(mutex);  
    readcount++;  
    if (readcount == 1) {  
        P(wrt);  
    }  
    V(mutex);  
    // reading  
    P(mutex);  
    readcount--;  
    if (readcount == 0) {  
        V(wrt);  
    }  
    V(mutex);  
}
```

Readers/Writers fixed

Writer:

```
while (1) {  
    P(wrt);  
    // writing  
    V(wrt);  
}
```

Reader:

```
while (1) {  
    P(mutex);  
    readcount++;  
    if (readcount == 1) {  
        P(wrt);  
    }  
    V(mutex);  
    // reading  
    P(mutex);  
    readcount--;  
    if (readcount == 0) {  
        V(wrt);  
    }  
    V(mutex);  
}
```

What if a Writer() is active, the first Reader() stalls on P(wrt), and additional Readers() try to enter?

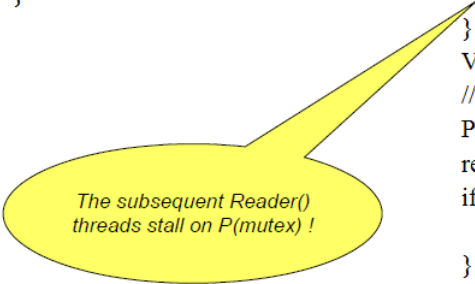
Readers/Writers fixed

Writer:

```
while (1) {  
    P(wrt);  
    // writing  
    V(wrt);  
}
```

Reader:

```
while (1) {  
    P(mutex);  
    readcount++;  
    if (readcount == 1) {  
        P(wrt);  
    }  
    V(mutex);  
    // reading  
    P(mutex);  
    readcount--;  
    if (readcount == 0) {  
        V(wrt);  
    }  
    V(mutex);  
}
```



The subsequent Reader() threads stall on P(mutex) !